



# Apache Spark

## In-Memory Data Processing

**September 2014 Meetup**

Organized by **Big Data Hyderabad Meetup Group**.

<http://www.meetup.com/Big-Data-Hyderabad/>

Rahul Jain

 @rahuldausa

# Agenda

- Why Spark
- Introduction
- Basics
- Hands-on
  - Installation
  - Examples

# Quick Questionnaire

How many people know/work on Scala ?


How many people know/work on Python ?

How many people know/heard/are using Spark ?


# Why Spark ?

- Most of Machine Learning Algorithms are iterative because each iteration can improve the results
- With Disk based approach each iteration's output is written to disk making it slow

## Hadoop execution flow



## Spark execution flow



<http://www.wiziq.com/blog/hype-around-apache-spark/>


# About Apache Spark



- Initially started at UC Berkeley in 2009
- Fast and general purpose cluster computing system
- 10x (on disk) - 100x (In-Memory) faster
- Most popular for running *Iterative Machine Learning Algorithms*.
- Provides high level APIs in
  - Java
  - Scala
  - Python
- Integration with Hadoop and its eco-system and can **read existing data**.
- <http://spark.apache.org/>


# Spark Stack

- Spark SQL
  - For SQL and unstructured data processing
- MLib
  - Machine Learning Algorithms
- GraphX
  - Graph Processing
- Spark Streaming
  - stream processing of live data streams



<http://spark.apache.org>

# Execution Flow



# Terminology

- **Application Jar**
  - User Program and its dependencies except Hadoop & Spark Jars bundled into a Jar file
- **Driver Program**
  - The process to start the execution (main() function)
- **Cluster Manager**
  - An external service to manage resources on the cluster (standalone manager, YARN, Apache Mesos)
- **Deploy Mode**
  - **cluster** : Driver inside the cluster
  - **client** : Driver outside of Cluster



# Terminology (contd.)

- **Worker Node** : Node that run the application program in cluster
- **Executor**
  - Process launched on a worker node, that runs the Tasks
  - Keep data in memory or disk storage
- **Task** : A unit of work that will be sent to executor
- **Job**
  - Consists multiple tasks
  - Created based on a Action
- **Stage** : Each Job is divided into smaller set of tasks called Stages that is sequential and depend on each other
- **SparkContext** :
  - represents the connection to a Spark cluster, and can be used to create RDDs, accumulators and broadcast variables on that cluster.

# Resilient Distributed Dataset (RDD)

- Resilient Distributed Dataset (RDD) is a basic Abstraction in Spark
- Immutable, Partitioned collection of elements that can be operated in parallel
- Basic Operations
  - map
  - filter
  - persist
- Multiple Implementation
  - [PairRDDFunctions](#) : RDD of Key-Value Pairs, groupByKey, Join
  - [DoubleRDDFunctions](#) : Operation related to double values
  - [SequenceFileRDDFunctions](#) : Operation related to SequenceFiles
- RDD main characteristics:
  - A list of partitions
  - A function for computing each split
  - A list of dependencies on other RDDs
  - Optionally, a Partitioner for key-value RDDs (e.g. to say that the RDD is hash-partitioned)
  - Optionally, a list of preferred locations to compute each split on (e.g. block locations for an HDFS file)
- Custom RDD can be also implemented (by overriding functions)

# Cluster Deployment

- Standalone Deploy Mode
  - simplest way to deploy Spark on a private cluster
- Amazon EC2
  - EC2 scripts are available
  - Very quick launching a new cluster
- Apache Mesos
- Hadoop YARN

# Monitoring

The screenshot shows a web browser window with the title "Spark shell - Spark Stages". The address bar shows "localhost:4040/stages/". The page has a navigation bar with the Spark logo and tabs for "Stages", "Storage", "Environment", and "Executors". The "Stages" tab is active. The main content area is titled "Spark Stages" and displays the following information:

- Total Duration: 45 s
- Scheduling Mode: FIFO
- Active Stages: 0
- Completed Stages: 0
- Failed Stages: 0

Below this information, there are three sections, each with a table header:

**Active Stages (0)**

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
----------	-------------	-----------	----------	------------------------	--------------	---------------

**Completed Stages (0)**

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
----------	-------------	-----------	----------	------------------------	--------------	---------------

**Failed Stages (0)**

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write	Failure Reason
----------	-------------	-----------	----------	------------------------	--------------	---------------	----------------

# Monitoring – Stages

Spark Pi - Spark Stages

localhost:4040/stages/

Spark Stages

Total Duration: 59 s  
Scheduling Mode: FIFO  
Active Stages: 1  
Completed Stages: 0  
Failed Stages: 0

Active Stages (1)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
0	(kill) reduce at SparkPi scala:35	2014/09/13 02:22:05	58 s	739/1000		

Completed Stages (0)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
----------	-------------	-----------	----------	------------------------	--------------	---------------

Failed Stages (0)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write	Failure Reason
----------	-------------	-----------	----------	------------------------	--------------	---------------	----------------

# Monitoring – Stages

Spark Pi - Details for Stage x

localhost4040/stages/stage/?id=0

Spark Stages Storage Environment Executors Spark Pi application UI

## Details for Stage 0

Total task time across all tasks: 41 s

### Summary Metrics for 73 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Result serialization time	0 ms	0 ms	0 ms	0 ms	21 ms
Duration	0.2 s	0.4 s	0.5 s	0.6 s	1.0 s
Time spent fetching task results	0 ms	0 ms	0 ms	0 ms	0 ms
Scheduler delay	5 ms	9 ms	14 ms	29 ms	1 s

### Aggregated Metrics by Executor

Executor ID	Address	Task Time	Total Tasks	Failed Tasks	Succeeded Tasks	Shuffle Read	Shuffle Write	Shuffle Spill (Memory)	Shuffle Spill (Disk)
localhost	CANNOT FIND ADDRESS	47 s	70	0	70	0.0 B	0.0 B	0.0 B	0.0 B

### Tasks

Task Index	Task ID	Status	Locality Level	Executor	Launch Time	Duration	GC Time	Result Ser Time	Errors
0	0	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.5 s			
1	1	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.3 s		21 ms	
2	2	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.8 s	98 ms	1 ms	
3	3	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.5 s			
4	4	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.7 s	98 ms		
5	5	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.7 s	98 ms		
6	6	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.7 s	98 ms	1 ms	
7	7	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.8 s	98 ms		
8	8	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:07	0.6 s	98 ms	1 ms	
9	9	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:07	0.3 s			

Let's try some examples...

# Spark Shell

```
./bin/spark-shell --master local[2]
```

The `--master` option specifies the [master URL for a distributed cluster](#), or `local` to run locally with one thread, or `local[N]` to run locally with N threads. You should start by using `local` for testing.

```
./bin/run-example SparkPi 10
```

This will run 10 iterations to calculate the value of Pi



# Basic operations...

```
scala> val textFile = sc.textFile("README.md")
textFile: spark.RDD[String] = spark.MappedRDD@2ee9b6e3
```

```
scala> textFile.count() // Number of items in this RDD
res0: Long = 126
```

```
scala> textFile.first() // First item in this RDD
res1: String = # Apache Spark
```

```
scala> val linesWithSpark = textFile.filter(line =>
line.contains("Spark"))
linesWithSpark: spark.RDD[String] = spark.FilteredRDD@7dd4af09
```

## Simplier - Single liner:

```
scala> textFile.filter(line => line.contains("Spark")).count()
// How many lines contain "Spark"?
res3: Long = 15
```

# Map - Reduce

```
scala> textFile.map(line => line.split(" ").size).reduce((a, b)
=> if (a > b) a else b)
res4: Long = 15
```

```
scala> import java.lang.Math
scala> textFile.map(line => line.split(" ").size).reduce((a, b)
=> Math.max(a, b))
res5: Int = 15
```

```
scala> val wordCounts = textFile.flatMap(line => line.split("
")).map(word => (word, 1)).reduceByKey((a, b) => a + b)
wordCounts: spark.RDD[(String, Int)] =
spark.ShuffledAggregatedRDD@71f027b8
wordCounts.collect()
```

# With Caching...

```
scala> linesWithSpark.cache()  
res7: spark.RDD[String] = spark.FilteredRDD@17e51082
```

```
scala> linesWithSpark.count()  
res8: Long = 15
```

```
scala> linesWithSpark.count()  
res9: Long = 15
```

# With HDFS...

```
val lines = spark.textFile("hdfs://...")  
val errors = lines.filter(line => line.startsWith("ERROR"))  
println(Total errors: + errors.count())
```

# Standalone (Scala)

```
/* SimpleApp.scala */
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._
import org.apache.spark.SparkConf

object SimpleApp {

  def main(args: Array[String]) {
    val logFile = "YOUR_SPARK_HOME/README.md" // Should be some file on your
system
    val conf = new SparkConf().setAppName("Simple Application")
.setMaster("local")
    val sc = new SparkContext(conf)
    val logData = sc.textFile(logFile, 2).cache()
    val numAs = logData.filter(line => line.contains("a")).count()
    val numBs = logData.filter(line => line.contains("b")).count()
    println("Lines with a: %s, Lines with b: %s".format(numAs, numBs))
  }
}
```

# Standalone (Java)

```
/* SimpleApp.java */
import org.apache.spark.api.java.*;
import org.apache.spark.SparkConf;
import org.apache.spark.api.java.function.Function;

public class SimpleApp {

    public static void main(String[] args) {

        String logFile = "YOUR_SPARK_HOME/README.md"; // Should be some file on your system
        SparkConf conf = new SparkConf().setAppName("Simple Application").setMaster("local");
        JavaSparkContext sc = new JavaSparkContext(conf);
        JavaRDD<String> logData = sc.textFile(logFile).cache();

        long numAs = logData.filter(new Function<String, Boolean>() {
            public Boolean call(String s) { return s.contains("a"); }
        }).count();

        long numBs = logData.filter(new Function<String, Boolean>() {
            public Boolean call(String s) { return s.contains("b"); }
        }).count();

        System.out.println("Lines with a: " + numAs + ", lines with b: " + numBs);
    }
}
```

# Standalone (Python)

```
"""SimpleApp.py"""
from pyspark import SparkContext

logFile = "YOUR_SPARK_HOME/README.md" # Should be some file on your
system
sc = SparkContext("local", "Simple App")
logData = sc.textFile(logFile).cache()

numAs = logData.filter(lambda s: 'a' in s).count()
numBs = logData.filter(lambda s: 'b' in s).count()

print "Lines with a: %i, lines with b: %i" % (numAs, numBs)
```

# Job Submission

```
$SPARK_HOME/bin/spark-submit \  
  --class "SimpleApp" \  
  --master local[4] \  
  target/scala-2.10/simple-project_2.10-1.0.jar
```



# Configuration

```
val conf = new SparkConf()  
    .setMaster("local")  
    .setAppName("CountingSheep")  
    .set("spark.executor.memory", "1g")  
  
val sc = new SparkContext(conf)
```

Questions ?

# Thanks!

@rahuldausa on twitter and slideshare  
<http://www.linkedin.com/in/rahuldausa>

**Join us @** For Solr, Lucene, Elasticsearch, Machine Learning, IR  
<http://www.meetup.com/Hyderabad-Apache-Solr-Lucene-Group/>  
<http://www.meetup.com/DataAnalyticsGroup/>

**Join us @** For Hadoop, Spark, Cascading, Scala, NoSQL, Crawlers and all cutting edge technologies.  
<http://www.meetup.com/Big-Data-Hyderabad/>